

Image Feature Extraction and Classification Using CNN Architectures

Dr. Meghna Kapoor

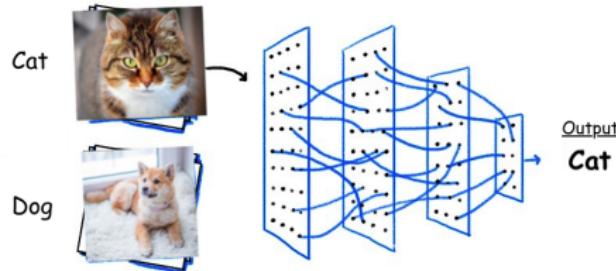
Learning Objective

- How features plays an important role in classification
- Implementation of different architectures for image classification

What is Image Classification

- Input is an image
- Output is a class label
- Example: cat, dog, car, person

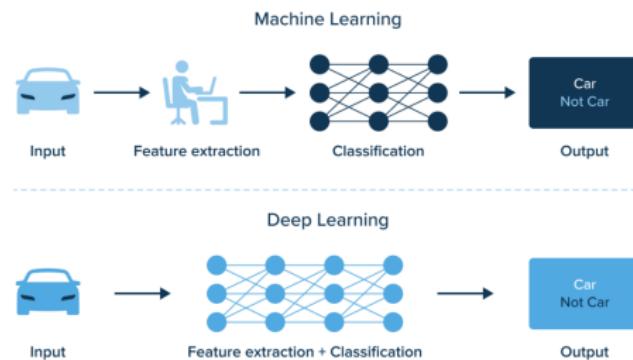
CNNs automatically learn useful patterns from images.



Why Do We Need CNNs

- Images have many pixels
- Manually designing features is difficult
- CNNs learn features directly from data

CNNs replace hand crafted features with learned features.



Two Main Parts of a CNN

1. Feature Extraction

- Convolution layers
- Pooling layers

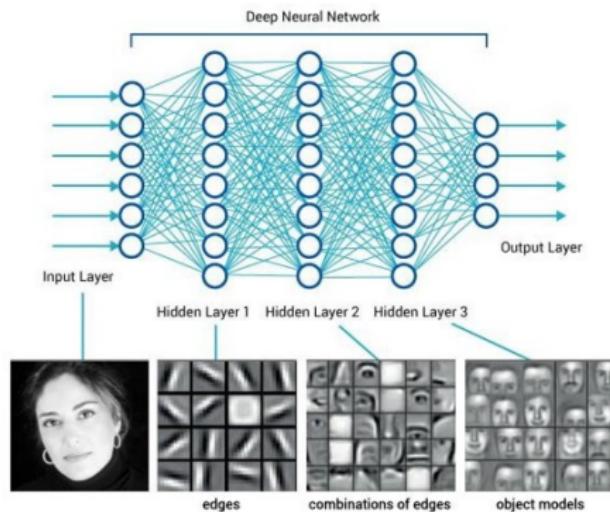
2. Classification

- Fully connected layers
- Softmax output

How Feature Extraction Works

- First layers detect edges
- Middle layers detect shapes
- Deeper layers detect objects

Feature learning becomes more meaningful with depth.



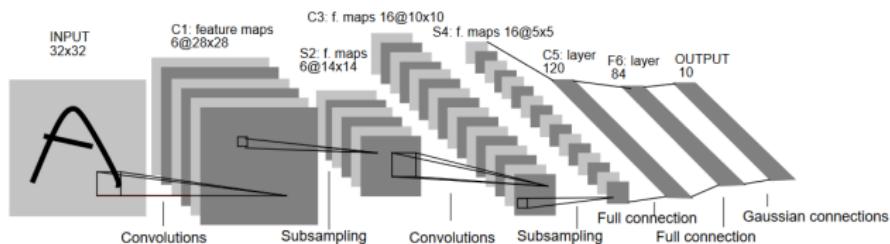
Why Count Parameters

- Parameters represent what the model learns
- More parameters means more memory usage
- More parameters increase training and inference time

Understanding parameter count helps choose the right model.

LeNet

- One of the first CNNs
- Few convolution layers
- Used for digit recognition



Simple model for simple images.

LeNet Parameters

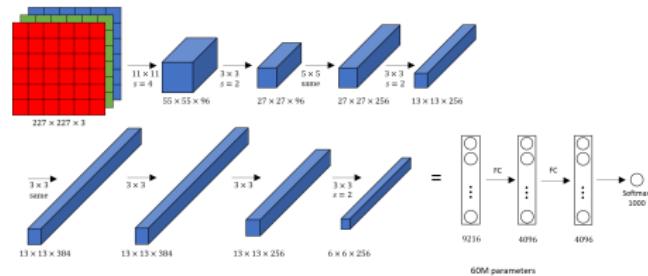
- Small filters
- Few channels
- Very low parameter count

Layer Name	Input W×H×D	Kernel W×H×D/S	Output W×H×D	Params	Mults
C1: conv2d	32×32×1	5×5×6	28×28×6	$1 \times 5 \times 5 \times 6 + 6 = 156$	$28 \times 28 \times 1 \times 5 \times 5 \times 6 = 117,600$
S2: pool/2	28×28×6	2×2/2	14×14×6	0	0
C3: conv2d	14×14×6	5×5×16	10×10×16	$6 \times 5 \times 5 \times 16 + 16 = 2,416$	$10 \times 10 \times 6 \times 5 \times 5 \times 16 = 240,000$
S4: pool/2	10×10×16	2×2/2	5×5×16	0	0
C5: conv2d	5×5×16	5×5×120	1×1×120	$16 \times 5 \times 5 \times 120 + 120 = 48,120$	$1 \times 1 \times 16 \times 5 \times 5 \times 120 = 48,000$
F6: conv2d	1×1×120	1×1×84	1×1×84	$120 \times 1 \times 1 \times 84 + 84 = 10,164$	$120 \times 84 = 10,080$
F7: conv2d	1×1×84	1×1×10	1×1×10	$84 \times 1 \times 1 \times 10 + 10 = 850$	$84 \times 40 = 840$
Total		61,706	416,520		

LeNet is lightweight and easy to train.

AlexNet

- Deeper than LeNet
- Works on larger images
- Uses ReLU and dropout



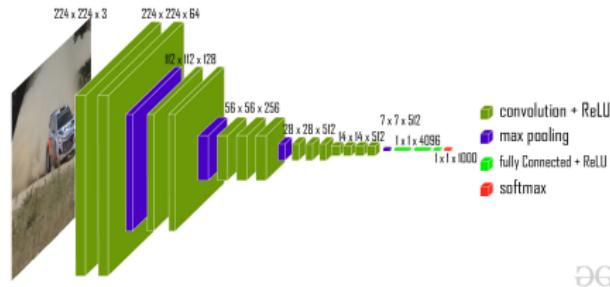
Showed that deep CNNs work well.

AlexNet Parameters

	Activation shape	Activation size	# parameters
Input image	227 x 227 x 3	154587	0
Conv 1	55 x 55 x 96 ($f=11 s = 4 p = 0$)	290400	34944
Pool 1	27 x 27 x 96 ($f=3 s = 2$)	69984	0
Conv 2	27 x 27 x 256 ($f=5 s = 1 p = 2$)	186624	614,656
Pool 2	13 x 13 x 256 ($f=3 s = 2$)	43264	0
Conv 3	13 x 13 x 384 ($f=3 s = 1 p = 1$)	64896	885,120
Conv 4	13 x 13 x 384 ($f=3 s = 1 p = 1$)	64896	1,327,488
Conv 5	13 x 13 x 256 ($f=3 s = 1 p = 1$)	43264	884,992
Pool 5	6 x 6 x 256 ($f=3 s = 2$)	9216	0
FC 3	4096 x 1	4096	37,748,737
FC 4	4096 x 1	4096	16,777,217
Softmax	1000 x 1	1000	4096001

VGGNet

- Many layers
- Uses only small 3×3 filters
- Easy to understand structure



∂G

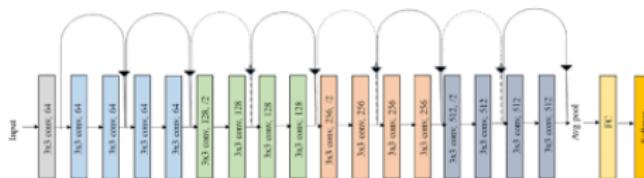
Deeper networks learn better features.

Problem with Very Deep Networks

- Training becomes difficult
- Gradients can vanish
- Accuracy may stop improving

ResNet

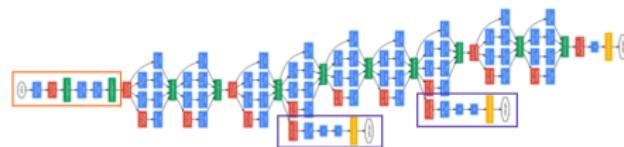
- Uses shortcut connections
- Skips some layers
- Makes training easier



Allows very deep CNNs.

InceptionNet

- Uses multiple filter sizes together
- Looks at image at different scales
- Efficient and accurate



GoogLeNet. The orange box is the stem, which has some preliminary convolutions. The purple boxes are auxiliary classifiers. The wide parts are the inception modules. (Source: Inception v1)

Combines multiple views of the image.