

# Image Feature Extraction and Classification Using CNN Architectures

Dr. Meghna Kapoor

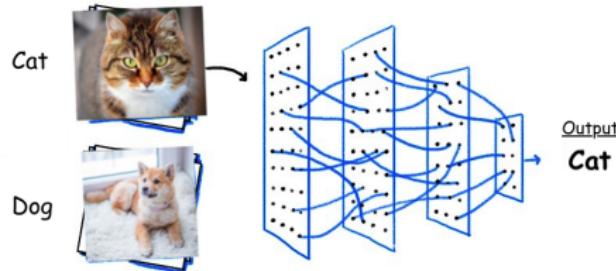
# Learning Objective

- How features plays an important role in classification
- Implementation of different architectures for image classification

# What is Image Classification

- Input is an image
- Output is a class label
- Example: cat, dog, car, person


CNNs automatically learn useful patterns from images.



# Why Do We Need CNNs

- Images have many pixels
- Manually designing features is difficult
- CNNs learn features directly from data

CNNs replace hand crafted features with learned features.



# Two Main Parts of a CNN

## 1. Feature Extraction

- Convolution layers
- Pooling layers


## 2. Classification

- Fully connected layers
- Softmax output

# How Feature Extraction Works

- First layers detect edges
- Middle layers detect shapes
- Deeper layers detect objects

**Feature learning becomes more meaningful with depth.**




# Why Count Parameters

- Parameters represent what the model learns
- More parameters means more memory usage
- More parameters increase training and inference time

Understanding parameter count helps choose the right model.

# LeNet

- One of the first CNNs
- Few convolution layers
- Used for digit recognition



Simple model for simple images.

# LeNet Parameters


- Small filters
- Few channels
- Very low parameter count

Layer Name	Input W×H×D	Kernel W×H×D/S	Output W×H×D	Params	Mults
C1: conv2d	32×32×1	5×5×6	28×28×6	$1 \times 5 \times 5 \times 6 + 6 = 156$	$28 \times 28 \times 1 \times 5 \times 5 \times 6 = 117,600$
S2: pool/2	28×28×6	2×2/2	14×14×6	0	0
C3: conv2d	14×14×6	5×5×16	10×10×16	$6 \times 5 \times 5 \times 16 + 16 = 2,416$	$10 \times 10 \times 6 \times 5 \times 5 \times 16 = 240,000$
S4: pool/2	10×10×16	2×2/2	5×5×16	0	0
C5: conv2d	5×5×16	5×5×120	1×1×120	$16 \times 5 \times 5 \times 120 + 120 = 48,120$	$1 \times 1 \times 16 \times 5 \times 5 \times 120 = 48,000$
F6: conv2d	1×1×120	1×1×84	1×1×84	$120 \times 1 \times 1 \times 84 + 84 = 10,164$	$120 \times 84 = 10,080$
F7: conv2d	1×1×84	1×1×10	1×1×10	$84 \times 1 \times 1 \times 10 + 10 = 850$	$84 \times 40 = 840$
Total		61,706		416,520	

LeNet is lightweight and easy to train.

# AlexNet

- Deeper than LeNet
- Works on larger images
- Uses ReLU and dropout




Showed that deep CNNs work well.

# AlexNet Parameters

	Activation shape	Activation size	# parameters
Input image	227 x 227 x 3	154587	0
Conv 1	55 x 55 x 96 ( $f=11 s = 4 p = 0$ )	290400	34944
Pool 1	27 x 27 x 96 ( $f=3 s = 2$ )	69984	0
Conv 2	27 x 27 x 256 ( $f=5 s = 1 p = 2$ )	186624	614,656
Pool 2	13 x 13 x 256 ( $f=3 s = 2$ )	43264	0
Conv 3	13 x 13 x 384 ( $f=3 s = 1 p = 1$ )	64896	885,120
Conv 4	13 x 13 x 384 ( $f=3 s = 1 p = 1$ )	64896	1,327,488
Conv 5	13 x 13 x 256 ( $f=3 s = 1 p = 1$ )	43264	884,992
Pool 5	6 x 6 x 256 ( $f=3 s = 2$ )	9216	0
FC 3	4096 x 1	4096	37,748,737
FC 4	4096 x 1	4096	16,777,217
Softmax	1000 x 1	1000	4096001

# VGGNet

- Many layers
- Uses only small  $3 \times 3$  filters
- Easy to understand structure



∂G


Deeper networks learn better features.

# Problem with Very Deep Networks

- Training becomes difficult
- Gradients can vanish
- Accuracy may stop improving

# ResNet


- Uses shortcut connections
- Skips some layers
- Makes training easier



Allows very deep CNNs.

# InceptionNet

- Uses multiple filter sizes together
- Looks at image at different scales
- Efficient and accurate




GoogLeNet. The orange box is the stem, which has some preliminary convolutions. The purple boxes are auxiliary classifiers. The wide parts are the inception modules. (Source: Inception v1)

Combines multiple views of the image.

# Dog Cat classification

- Binary image classification problem
- Real-world dataset with large visual variability
- More complex than digit recognition
- Tests generalization capability of CNNs



# Problem Statement

- Input: RGB images of dogs and cats
- Output: Class label (Dog or Cat)
- Dataset organized using folder structure
- Supervised learning setup

# Dataset Organization

- Images stored in Folder
- Two folders: Train and Test
- Each folder contains two classes
- Used ImageFolder for automatic labeling

# CNN Architectures Used

- LeNet-5
- AlexNet
- VGG16
- ResNet18
- Inception Network

# Why Compare Multiple Architectures

- Different depth and complexity
- Different feature learning capacity
- Different computational cost
- Observe performance trends

# Preprocessing and Augmentation

- Images resized to required input size
- Normalization applied
- Data augmentation used only during training
- Improves robustness and generalization

# Training Setup

- Loss function: Cross-Entropy Loss
- Optimizer: Adam
- Batch size fixed across models
- Same train-test split for fair comparison

# Evaluation Metrics: Mathematical Definitions

Let:

- $TP$  = True Positives
- $TN$  = True Negatives
- $FP$  = False Positives
- $FN$  = False Negatives
- **Accuracy**

$$\text{Accuracy} = \frac{TP + TN}{TP + TN + FP + FN}$$

- **Precision**

$$\text{Precision} = \frac{TP}{TP + FP}$$

- **Recall**

$$\text{Recall} = \frac{TP}{TP + FN}$$

- **F1-score**

$$\text{F1-score} = 2 \times \frac{\text{Precision} \times \text{Recall}}{\text{Precision} + \text{Recall}}$$

# Accuracy Comparison

- LeNet shows lowest accuracy
- AlexNet improves feature learning
- VGG16 benefits from deeper structure
- ResNet18 converges faster
- Inception achieves strong generalization


# Confusion Matrix Analysis

- Shows class-wise prediction errors
- Highlights false positives and false negatives
- Helps understand model bias
- Deeper models reduce misclassification

		True Class	
		Positive	Negative
Predicted Class	Positive	TP	FP
	Negative	FN	TN

# ROC Curve and AUC

- Threshold independent evaluation
- Measures discrimination ability
- Higher AUC indicates stronger classifier
- Inception and ResNet show higher AUC



# Precision, Recall, and F1 Score

- Precision measures prediction reliability
- Recall measures detection capability
- F1 score balances both
- Important when class imbalance exists

# Qualitative Results

- Visual inspection of predictions
- True label vs predicted label displayed
- Confidence score indicates certainty
- Misclassifications analyzed visually

# Key Observations

- Deeper models learn richer features
- Residual connections improve stability
- Inception captures multi-scale patterns
- Model complexity improves performance

# Limitations

- Training time increases with depth
- Large models require more memory
- Small dataset may cause overfitting

# Conclusion

- CNNs outperform traditional methods
- Architecture choice impacts accuracy
- Deeper and multi-scale models perform better